

Service

In Kubernetes, a Service is a method for exposing a network application that is running as one or more Pods in your cluster.

Types:

- 1)clusterIp
- 2)Nodeport
- 3)Load Balancer
- 4)Headless
- 5)External name service
- 6) External IP service

What are Services in Kubernetes?

Think of a service as a stable, consistent way to access your applications running inside a Kubernetes cluster. It's like a virtual front door that sits in front of your Pods, making them accessible without needing to know the specific details of where they're running.

Key Features of Services:

Abstraction: Services abstract away the details of where your Pods are running. You don't need to know the specific IP addresses or ports of your Pods to access them.

Load Balancing: Services can distribute traffic across multiple Pods, ensuring that your application is scalable and resilient.

Service Discovery: Services provide a way for Pods to find each other within the cluster.

External Access: Services can be exposed to the outside world, allowing users to access your application from outside the cluster.

How Services Work:

Service Definition: You define a Service using a YAML or JSON file. This file specifies the name of the service, the selector (which Pods it should route traffic to), and the port mappings.

Service Discovery: Kubernetes automatically creates a service endpoint for the service. This endpoint is a virtual IP address and port that can be used to access the service.

Traffic Routing: When a request is made to the service endpoint, Kubernetes routes the request to one of the Pods that match the service's selector. Kubernetes uses a load balancer to distribute traffic across the Pods.

```
apiVersion: v1
kind: Service
metadata:
  name: my-web-app
spec:
  selector:
    app: my-web-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

This Service definition specifies:

Name: my-web-app

Selector: It will route traffic to Pods with the label app : my-web-app.

Port Mapping: It maps port 80 on the service to port 8080 on the Pods.

- 1) **ClusterIP:** This service will be accessible only within the Kubernetes cluster. It will not be exposed to the outside world.

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
  type: ClusterIP
```

Explanation:

apiVersion: Specifies the Kubernetes API version (v1 in this case).

kind: Indicates that this is a Service definition.

metadata:

name: The name of the service. Here, it's "my-service".

spec: Defines the service's configuration.

selector: A label selector that matches the Pods the service should route traffic to. In this case, it will route traffic to Pods with the label `app : my-app`.

ports: An array of port mappings.

protocol: The protocol used (TCP in this case).

port: The port that the service listens on. Here, it's port 80.

targetPort: The port that the Pods listen on. Here, it's port 8080.

type: The type of service. Here, it's set to "**ClusterIP**", which means the service will be accessible only within the Kubernetes cluster.

ClusterIP: This service will be accessible only within the Kubernetes cluster. It will not be exposed to the outside world.

Selector: The service will route traffic to Pods that have the label `app : my-app`.

Port Mapping: The service maps port 80 to port 8080 on the Pods.

To Create the Service:

Save the YAML: Save this YAML code to a file named `my-service.yaml`.

Apply the YAML: Run the following command to create the service:

```
kubectl apply -f my-service.yaml
```

Kubernetes will then create the service and make it available within the cluster.

Important:

You'll need to have Pods running with the label `app : my-app` for this service to work correctly.

To access the service from within the cluster, you can use the service's name and port. For example, you could use `curl my-service:80` to access the service.

What is a NodePort Service?

Think of a NodePort service as a way to make your application accessible from outside your Kubernetes cluster by assigning a specific port on each node in the cluster. It's like opening a window on each house in your neighborhood, allowing people to access your application from the street.

Key Features of NodePort Services:

External Access: NodePort services allow you to access your application from outside the Kubernetes cluster.

Simple Configuration: They're relatively easy to configure, making them a good choice for simple deployments.

Port Mapping: Each node in the cluster will have a specific port assigned to the service. This port will be mapped to the port that your application is listening on inside the Pods.

How NodePort Services Work:

Service Definition: You define a NodePort service using a YAML or JSON file. This file specifies the name of the service, the selector (which Pods it should route traffic to), and the port mappings.

Port Assignment: Kubernetes automatically assigns a unique port on each node in the cluster to the service.

Traffic Routing: When a request is made to the assigned port on a node, Kubernetes routes the request to one of the Pods that match the service's selector. Kubernetes uses a load balancer to distribute traffic across the Pods.

```
apiVersion: v1
kind: Service
metadata:
  name: my-web-app
spec:
  selector:
    app: my-web-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
      nodePort: 30080 # Assign port 30080 on each node
  type: NodePort
```

Explanation:

apiVersion: Specifies the Kubernetes API version (v1 in this case).

kind: Indicates that this is a Service definition.

metadata:

name: The name of the service. Here, it's "my-web-app".

spec: Defines the service's configuration.

selector: A label selector that matches the Pods the service should route traffic to. In this case, it will route traffic to Pods with the label `app: my-app`.

ports: An array of port mappings.

protocol: The protocol used (TCP in this case).

port: The port that the service listens on. Here, it's port 80.

targetPort: The port that the Pods listen on. Here, it's port 8080.

nodePort: The port that will be assigned on each node. Here, it's port 30080.

type: The type of service. Here, it's set to "**NodePort**", which means the service will be accessible from outside the cluster.

To Create the Service:

Save the YAML: Save this YAML code to a file named `my-web-app.yaml`.

Apply the YAML: Run the following command to create the service:

```
kubectl apply -f my-web-app.yaml
```

Kubernetes will then create the service and assign a unique port on each node.

Accessing the Service:

From Inside the Cluster: You can access the service using its name and port (e.g., `curl my-web-app:80`).

From Outside the Cluster: You can access the service by using the assigned NodePort on each node. For example, if your node's IP address is `10.128.0.10`, you could access the service using `curl 10.128.0.10:30080`.

Important:

You'll need to have Pods running with the label `app: my-web-app` for this service to work correctly.

The assigned NodePort will be in the range of 30000-32767. Kubernetes will automatically choose a port within this range.

What is a LoadBalancer Service?

Think of a LoadBalancer service as a way to make your application accessible from the internet using a load balancer. It's like having a dedicated receptionist who handles all incoming calls and directs them to the appropriate person (Pod) within your organization (Kubernetes cluster).

Key Features of LoadBalancer Services:

External Access: LoadBalancer services allow you to access your application from the internet.

High Availability: The load balancer distributes traffic across multiple Pods, ensuring that your application remains available even if some Pods fail.

Scalability: The load balancer can handle a large volume of traffic, making your application scalable.

Security: Load balancers can provide security features like SSL termination and access control.

How LoadBalancer Services Work:

Service Definition: You define a LoadBalancer service using a YAML or JSON file. This file specifies the name of the service, the selector (which Pods it should route traffic to), and the port mappings.

Load Balancer Creation: When you create a LoadBalancer service, Kubernetes will automatically create a load balancer in your cloud provider's infrastructure.

Traffic Routing: When a request is made to the load balancer's IP address, the load balancer will route the request to one of the Pods that match the service's selector. The load balancer uses a load balancing algorithm to distribute traffic across the Pods.

```
apiVersion: v1
kind: Service
metadata:
  name: my-web-app
spec:
  selector:
    app: my-web-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
  type: LoadBalancer
```

Explanation:

apiVersion: Specifies the Kubernetes API version (v1 in this case).

kind: Indicates that this is a Service definition.

metadata:

name: The name of the service. Here, it's "my-web-app".

spec: Defines the service's configuration.

selector: A label selector that matches the Pods the service should route traffic to. In this case, it will route traffic to Pods with the label `app: my-web-app`.

ports: An array of port mappings.

protocol: The protocol used (TCP in this case).

port: The port that the service listens on. Here, it's port 80.

targetPort: The port that the Pods listen on. Here, it's port 8080.

type: The type of service. Here, it's set to "**LoadBalancer**", which means the service will be exposed to the internet using a load balancer.

To Create the Service:

Save the YAML: Save this YAML code to a file named `my-web-app.yaml`.

Apply the YAML: Run the following command to create the service:

```
kubectl apply -f my-web-app.yaml
```

Accessing the Service:

From the Internet: You can access the service using the load balancer's public IP address. The IP address will be displayed in the `LoadBalancer Ingress` field when you run `kubectl get services my-web-app`.

Important:

You'll need to have Pods running with the label `app: my-web-app` for this service to work correctly.

The load balancer's IP address may take a few minutes to become available. You can check the status of the service using `kubectl get services my-web-app`.

What is a Headless Service?

Think of a headless service as a way to group Pods together without providing a central point of access. It's like having a list of phone numbers for all the people in your organization, but no single receptionist to answer calls. Instead, you have to call each person directly.

Key Features of Headless Services:

No Cluster IP: Headless services don't have a cluster IP address. Instead, they expose the IP addresses and ports of the Pods directly.

Service Discovery: Headless services are primarily used for service discovery within the cluster. Pods can use the service's DNS name to find the IP addresses and ports of other Pods in the service.

Direct Communication: Pods can communicate with each other directly using the IP addresses and ports exposed by the headless service.

How Headless Services Work:

Service Definition: You define a headless service using a YAML or JSON file. This file specifies the name of the service, the selector (which Pods it should route traffic to), and the port mappings.

DNS Records: Kubernetes automatically creates DNS records for the headless service. These records map the service's name to the IP addresses and ports of the Pods in the service.

Service Discovery: Pods can use the service's DNS name to resolve the IP addresses and ports of other Pods in the service.

```
apiVersion: v1
kind: Service
metadata:
  name: my-headless-service
```

```
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
  clusterIP: None # This is what makes it headless
```

Explanation:

apiVersion: Specifies the Kubernetes API version (v1 in this case).

kind: Indicates that this is a Service definition.

metadata:

name: The name of the service. Here, it's "my-headless-service".

spec: Defines the service's configuration.

selector: A label selector that matches the Pods the service should route traffic to. In this case, it will route traffic to Pods with the label `app: my-app`.

ports: An array of port mappings.

protocol: The protocol used (TCP in this case).

port: The port that the service listens on. Here, it's port 80.

targetPort: The port that the Pods listen on. Here, it's port 8080.

clusterIP: This is set to "**None**", which makes the service headless.

To Create the Service:

Save the YAML: Save this YAML code to a file named `my-headless-service.yaml`.

Apply the YAML: Run the following command to create the service:

```
kubectl apply -f my-headless-service.yaml
```

Kubernetes will then create the headless service and create DNS records for it.

Accessing the Service:

From Pods: Pods can use the service's DNS name to resolve the IP addresses and ports of other Pods in the service. For example, a Pod could use `my-headless-service.default.svc.cluster.local:80` to access another Pod in the service.

Important:

You'll need to have Pods running with the label `app: my-app` for this service to work correctly.

Headless services are typically used for service discovery within the cluster. They are not intended for external access.

Beyond the Basics: More Service Types and Concepts

ExternalName Service: This type of service is used to expose an external service or a DNS name. It's useful for situations where you want to access a service that's not running within your Kubernetes cluster.

"ExternalName" is a type of service in Kubernetes, but it's not about assigning an external IP address directly. It's more about creating a DNS alias within your cluster that points to an external service or a DNS name.

Let's break it down:

ExternalName Service: A DNS Alias for External Services

Purpose: ExternalName services are used to expose an external service or a DNS name within your Kubernetes cluster. This means you can access an external service using a DNS name that's defined within your cluster.

No Cluster IP: ExternalName services don't have a cluster IP address. They don't manage Pods or route traffic within the cluster.

DNS Mapping: The key feature of ExternalName services is that they create a DNS mapping within your cluster. When you create an ExternalName service, you specify an `externalName` field, which is the DNS name or IP address of the external service you want to expose. Kubernetes will then create a DNS record that maps the service's name within your cluster to the specified `externalName`.

```
apiVersion: v1
kind: Service
metadata:
  name: my-external-service
spec:
  type: ExternalName
  externalName: example.com
```

type: ExternalName: This specifies that this is an ExternalName service.

externalName: example.com: This defines the external DNS name or IP address that the service will map to.

How it Works:

Service Creation: You create the ExternalName service using the YAML file above.

DNS Mapping: Kubernetes creates a DNS record that maps the service's name (`my-external-service.default.svc.cluster.local`) to the specified `externalName` (`example.com`).

Access: Pods within your cluster can now access the external service using the service's name.

For example, a Pod could use `curl my-external-service.default.svc.cluster.local` to access the service at `example.com`.

Key Points:

No Traffic Routing: ExternalName services don't route traffic within the cluster. They simply create a DNS mapping.

External Access: They are used to expose external services or DNS names within your cluster.

Simple Configuration: They are relatively easy to configure.

External IPs

If there are external IPs that route to one or more cluster nodes, Kubernetes Services can be exposed on those `externalIPs`. When network traffic arrives into the cluster, with the external IP (as destination IP) and the port matching that Service, rules and routes that Kubernetes has configured ensure that the traffic is routed to one of the endpoints for that Service.

When you define a Service, you can specify `externalIPs` for any [service type](#)

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app.kubernetes.io/name: MyApp
  ports:
    - name: http
```

```
protocol: TCP
port: 80
targetPort: 49152
externalIPs:
- 198.51.100.32
```

Explanation:

apiVersion: Specifies the Kubernetes API version (v1 in this case).

kind: Indicates that this is a Service definition.

metadata:

name: The name of the service. Here, it's "my-service".

spec: Defines the service's configuration.

selector: A label selector that matches the Pods the service should route traffic to. In this case, it will route traffic to Pods with the label `app.kubernetes.io/name: MyApp`.

ports: An array of port mappings.

name: The name of the port. Here, it's "http".

protocol: The protocol used (TCP in this case).

port: The port that the service listens on. Here, it's port 80.

targetPort: The port that the Pods listen on. Here, it's port 49152.

externalIPs: An array of external IP addresses that will be assigned to the service. Here, it's `198.51.100.32`.

Key Points:

External IP: This service will be accessible from outside the Kubernetes cluster using the external IP address `198.51.100.32`.

Selector: The service will route traffic to Pods that have the label `app.kubernetes.io/name: MyApp`.

Port Mapping: The service maps port 80 to port 49152 on the Pods.

How it Works:

Service Creation: You create the service using this JSON file.

External IP Assignment: Kubernetes will assign the external IP address `198.51.100.32` to the service.

Traffic Routing: When a request is made to the external IP address, Kubernetes will route the request to one of the Pods that match the service's selector. Kubernetes uses a load balancer to distribute traffic across the Pods.

Important Considerations:

External IP Ownership: You need to ensure that you have control over the external IP address 198.51.100.32. This IP address should be assigned to your Kubernetes cluster or your cloud provider's infrastructure.

Cloud Provider Support: The ability to assign external IP addresses to services may vary depending on your cloud provider. Check your cloud provider's documentation for specific instructions.

To Create the Service:

Save the JSON: Save this JSON code to a file named `my-service.json`.

Apply the JSON: Run the following command to create the service:

```
kubectl apply -f my-service.yaml
```

Kubernetes will then create the service and assign the external IP address to it.

Accessing the Service:

From the Internet: You can access the service using the external IP address 198.51.100.32.

Use Cases for Different Service Types:

ClusterIP: Ideal for internal services that are only accessible within the Kubernetes cluster.

NodePort: Useful for exposing services to the outside world in a simple and straightforward manner.

LoadBalancer: The best choice for public-facing applications that require high availability, scalability, and security.

ExternalName: Used to expose external services or DNS names.

Headless: Primarily used for service discovery within the cluster.

Key Points:

Choose the Right Service Type: The type of service you choose will depend on your specific needs and requirements.

Service Discovery is Crucial: Service discovery is essential for enabling communication between services in your Kubernetes cluster.

Service Meshes Provide Advanced Features: Service meshes can provide advanced features for managing and securing communication between services.

Working on Scheduling the Pod to Node will come back with Document by . Happy Learning 😊....

